

hw1

September 12, 2017

1 CSCI 5992 - HW1 - Kelly/Milne

1.1 Part 1

1.1.1 Solution

1a) $w_1 = -2.04427331$ $w_2 = 3.99683416$ $b = -0.92427055$

1b) Levenberg–Marquardt as per implementation below

```
In [1]: ###
import os
import numpy as np
import scipy.io
import scipy.optimize as optimization
#from IPython.core.debugger import set_trace

# Data
filedir = os.path.dirname(os.path.realpath('__file__'))
datapath = os.path.join(filedir, 'assign1_data.mat')
data = scipy.io.loadmat(datapath)

# fitter wants shape (k,M) (k number of predictors)
# data['x'] has shape (M,k) so we fix that
x1 = data['x'][:,0]
x2 = data['x'][:,1]
X = np.array([x1,x2])
# argg - shape (100,1) - gives obtuse error in the fitter
# we fix that too (flatten)
y = np.array(data['y']).flatten()

# Initial guess
w0 = np.array([1, 1, 1])

# Objective function
# y = w1 * x1 + w2 * x2 + b.
def func(X, b, w1, w2):
    # unpack independent vars
    rval = w1*X[0] + w2*X[1] + b
```

```

    return rval

result = optimization.curve_fit(func, X, y, w0)
print(result[0])

[-0.92427055 -2.04427331  3.99683416]

```

1.2 Part 2

1.2.1 Solution

2a) $w_1 = -2.0427036$ $w_2 = 3.99155586$ $b = -0.92233149$

2b) Initially we batch processed using the full 100 samples. A step size of 0.05 diverged so used 0.01. Approx $N=100$ epochs to recover the results from the LM method. Termination was based on plot below of error vs epoch, which amounts to a % change in the error function of $\sim 10^{-4}$ between epochs.

```

In [2]: import os
import numpy as np
import scipy.io
import scipy.optimize as optimization
#from IPython.core.debugger import set_trace

EPS = 0.01
N_EPOCH = 100
err = []
def dw(w,d,X):
    # using numpy einstein summation to vectorize the computation
    # vector of coefficients (d-w.x) for each pattern (alpha)
    C = d - np.einsum('i,ij->j', w, X) # shape is (100,)
    e = 1.0/len(C)*np.sum([ c**2 for c in C])
    err.append(e)
    # C_i * X_ji or C * transpose(X)
    dw = EPS*np.einsum('i,ji->j', C, X) # shape is (3,)
    return dw

def main():
    # data
    filedir = os.path.dirname(os.path.realpath('__file__'))
    datapath = os.path.join(filedir, 'assign1_data.mat')
    data = scipy.io.loadmat(datapath)
    x1 = data['x'][:,0]
    x2 = data['x'][:,1]
    # inputs - with bias input tied high, shape is (3,100)
    X = np.array([np.full(len(x1),1),x1,x2])
    # output
    d = np.array(data['y']).flatten()
    # starting vector of weights

```

```

w = [1,1,1]

for _ in range(N_EPOCH):
    w = w + dw(w,d,X)
print(w)

# invoke main
main()

[-0.92263116 -2.04308226  3.99249859]


```

```

In [3]: # To plot pretty figures
%matplotlib inline
import matplotlib
import matplotlib.pyplot as plt
plt.rcParams['axes.labelsize'] = 14
plt.rcParams['xtick.labelsize'] = 12
plt.rcParams['ytick.labelsize'] = 12
ax = plt.subplot(111)
plt.semilogx(err)
ax.set_ylim([0, 2])
plt.ylabel('Error')
plt.xlabel('Epoch')
plt.title('Error vs Epoch for Gradient Descent Linear Regression')

Out[3]: <matplotlib.text.Text at 0x115214278>

```



We experimented with minibatch size of 25 which converged more smoothly. See attached mini_batch.png

1.3 Part 3

1.3.1 Solution

3a) $b = -16$. $w_1 = -50.43661604$ $w_2 = 81.86640442$

3b) See plot in 'Visualization of Perceptron Results'

```
In [4]: import os as os
import numpy as np
import scipy.io
import scipy.optimize as optimization
from IPython.core.debugger import set_trace

N_EPOCH = 30

n_mis = []
def dwp(w,d,X):
    # using numpy einstein summation to vectorize the computation
    # calculating w.x.d -- w.x.d < 0 => incorrect classification
    C = d*np.einsum('i,ij->j', w, X) # shape is (100,)
    XT = X.transpose()
    # use enumerate for the equivalent to each_with_index (Ruby)
    xd = [ d[i]*XT[i] for i,c in enumerate(C) if c < 0]
    # for N vs EPOCH plot
    n_mis.append(len(xd))
    # sum xd element-wise
    # dw is x.t summed over misclassified teachers
    dw = np.einsum('ij->j',xd) if(len(xd) > 0) else 0
    return dw

def main():
    # data
    filedir = os.path.dirname(os.path.realpath('__file__'))
    datapath = os.path.join(filedir,'assign1_data.mat')
    data = scipy.io.loadmat(datapath)
    x1 = data['x'][:,0]
    x2 = data['x'][:,1]
    # inputs - with bias input tied high, shape is (3,100)
    X = np.array([np.full(len(x1),1),x1,x2])
    # output
    d = np.array(data['z']).flatten()
    # use domain -1,1 for teachers so we can use x.w.d<0 for the classification test
    d = [z if z == 1 else -1 for z in d]
    # starting vector of weights
```

```

w = [1,1,1]
# sweep through the data N_EPOCH times
for _ in range(N_EPOCH):
    w = w + dwp(w,d,X)

return w

# invoke main
w_perceptron = main()
print(w_perceptron)

[-16.          -50.43661604   81.86640442]

```

2 Visualization of perceptron results

Read the data into a panda dataframe for visualization

```

In [5]: # Common imports
import numpy as np
import os

# to make this notebook's output stable across runs
np.random.seed(42)

# To plot pretty figures
%matplotlib inline
import matplotlib
import matplotlib.pyplot as plt
plt.rcParams['axes.labelsize'] = 14
plt.rcParams['xtick.labelsize'] = 12
plt.rcParams['ytick.labelsize'] = 12

# Where to save the figures
PROJECT_ROOT_DIR = "."

import pandas as pd

def load_data():
    csv_path = os.path.join(".", "assign1_data.txt")
    return pd.read_csv(csv_path, sep="\s+")

def load_data():
    csv_path = os.path.join(".", "assign1_data.txt")
    return pd.read_csv(csv_path, sep="\s+")

data = load_data()
data.head()

```

```
c0 = data[(data.z==0)]
c1 = data[(data.z==1)]
```

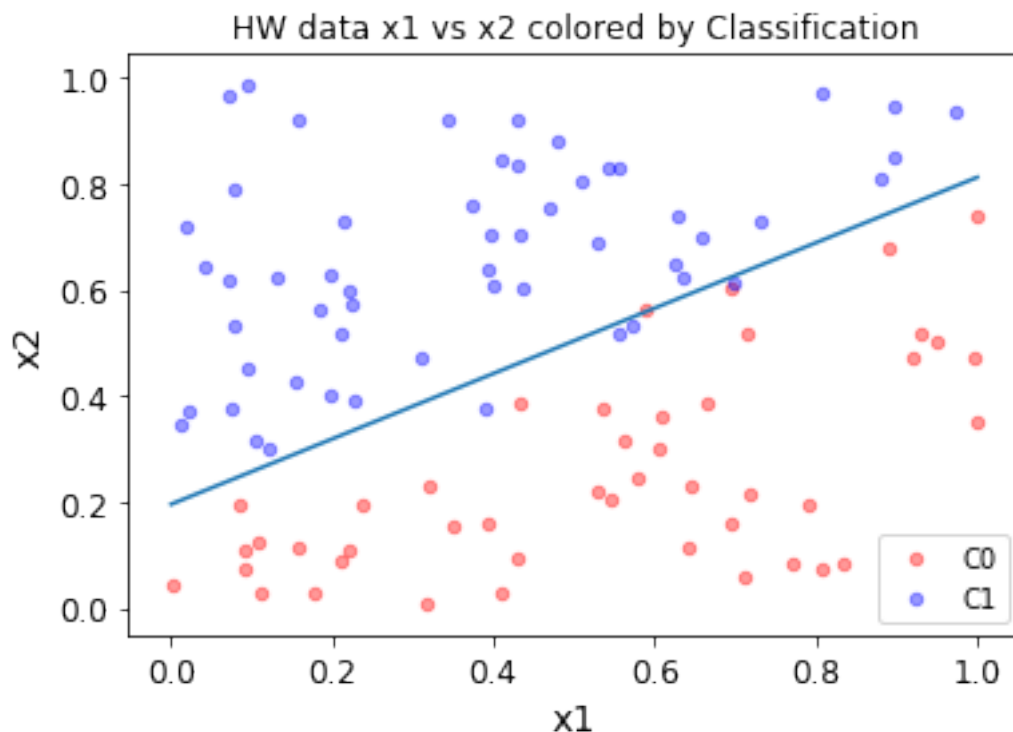
2.0.1 HW data in parameter space with classification indicated by color

The plotted line is the surface of separation given by $x_2 = -1/w_2(w_1x_1 + w_0)$ where w is the weight vector calculated using perceptron algorithm. As is indicated in the figure the data is not linearly separable so the perceptron algorithm does not converge, instead it jitters around near the minimum.

```
In [6]: X = np.linspace(0, 1, 128, endpoint=True)
        W = w_perceptron
        L = -(1/W[2])*(W[1]*X + W[0])


        ax = c0.plot.scatter(x="x1", y="x2", color='Red', alpha=0.4, label='C0')
        c1.plot.scatter( x="x1", y="x2", color='Blue', alpha=0.4, label='C1', ax=ax)
        ax.plot(X,L)
        plt.title('HW data x1 vs x2 colored by Classification')
```

```
Out[6]: <matplotlib.text.Text at 0x1169c1a58>
```



```
In [7]: plt.ylabel(r'$N_{\text{misclassified}}$')
        plt.xlabel('Epoch')
        plt.title('Number of Misclassifications vs Epoch')
        plt.plot(n_mis)
```

Out [7]: [<matplotlib.lines.Line2D at 0x116ae3c18>]



3 Part 4

Plot of performance vs number of training samples below.

```
In [8]: import os as os
import numpy as np
import scipy.io
import scipy.optimize as optimization

N_EPOCH = 20

n_mis = []
def dwp(w,d,X):
    # using numpy einsum summation to vectorize the computation
    # calculating w.x.d -- w.x.d < 0 => incorrect classification
    C = d*np.einsum('i,ij->j', w, X) # shape is (100,)
    XT = np.einsum('ij->ji', X)
    # use enumerate for the equivalent to each_with_index (Ruby)
    xd = [ d[i]*XT[i] for i,c in enumerate(C) if c < 0]
    # sum xd element-wise
    # dw is x.t summed over misclassified teachers
```

```

dw = np.einsum('ij->j',xd) if(len(xd) > 0) else 0 # shape is (3,)
return dw

def train(n, X, d):
    # use n_sample for training
    X = X[:,0:n]
    d = d[0:n]
    # starting vector of weights
    w = [1,1,1]
    # sweep through the data N_EPOCH times
    for _ in range(N_EPOCH):
        w = w + dwp(w,d,X)

    return w

def performance(w,X,d):
    X = X[:,75:]
    d = d[75:]
    C = d*np.einsum('i,ij->j', w, X) # shape is (100,)
    XT = np.einsum('ij->ji', X)
    # d.x.t for misclassified
    xd = [ d[i]*XT[i] for i,c in enumerate(C) if c < 0]
    return len(xd)/25.0

# invoke main
def main():
    # data
    filedir = os.path.dirname(os.path.realpath('__file__'))
    datapath = os.path.join(filedir,'assign1_data.mat')
    data = scipy.io.loadmat(datapath)
    x1 = data['x'][:,0]
    x2 = data['x'][:,1]
    # inputs - with bias input tied high, shape is (3,100)
    X = np.array([np.full(len(x1),1),x1,x2])
    # output
    d = np.array(data['z']).flatten()
    # use domain -1,1 for teachers so we can use x.w.d<0 for the classification test
    d = [z if z == 1 else -1 for z in d]
    samples = [5, 10, 25, 50, 75]
    perf = []
    for n in samples:
        w = train(n, X, d)
        p = performance(w, X, d)
        perf.append(1-p)
    return samples, perf

samples, perf = main()

```



```
In [9]: plt.ylabel('% Correct')
plt.xlabel(r'$N_{\text{training}}$')
plt.title(r'Percent Correctly Classified vs Number Training Samples')
plt.plot(samples, perf)
```

```
Out[9]: [<matplotlib.lines.Line2D at 0x116b02c50>]
```

